

R has many `*apply` functions which are ably described in the help files (e.g. `?apply`). There are enough of them, though, that beginning users may have difficulty deciding which one is appropriate for their situation or even remembering them all. They may have a general sense that "I should be using an `*apply` function here", but it can be tough to keep them all straight at first.

Despite the fact (noted in other answers) that much of the functionality of the `*apply` family is covered by the extremely popular `plyr` package, the base functions remain useful and worth knowing.

This answer is intended to act as a sort of **signpost** for new users to help direct them to the correct `*apply` function for their particular problem. Note, this is **not** intended to simply regurgitate or replace the R documentation! The hope is that this answer helps you to decide which `*apply` function suits your situation and then it is up to you to research it further. With one exception, performance differences will not be addressed.

- **apply** - *When you want to apply a function to the rows or columns of a matrix (and higher-dimensional analogues).*
- # Two dimensional matrix
- `M <- matrix(seq(1,16), 4, 4)`
- 
- # apply min to rows
- `apply(M, 1, min)`
- `[1] 1 2 3 4`
- 
- # apply max to columns
- `apply(M, 2, max)`
- `[1] 4 8 12 16`
- 
- # 3 dimensional array
- `M <- array( seq(32), dim = c(4,4,2))`
- 
- # Apply sum across each `M[, , ]` - i.e Sum across 2nd and 3rd dimension
- `apply(M, 1, sum)`
- # Result is one-dimensional
- `[1] 120 128 136 144`
- 
- # Apply sum across each `M[, *, ]` - i.e Sum across 3rd dimension
- `apply(M, c(1,2), sum)`
- # Result is two-dimensional
- |      | [,1] | [,2] | [,3] | [,4] |
|------|------|------|------|------|
| [1,] | 18   | 26   | 34   | 42   |
| [2,] | 20   | 28   | 36   | 44   |
| [3,] | 22   | 30   | 38   | 46   |
| [4,] | 24   | 32   | 40   | 48   |

If you want row/column means or sums for a 2D matrix, be sure to investigate the highly optimized, lightning-quick `colMeans`, `rowMeans`, `colSums`, `rowSums`.

- **lapply** - *When you want to apply a function to each element of a list in turn and get a list back.*

This is the workhorse of many of the other \*apply functions. Peel back their code and you will often find `lapply` underneath.

```
x <- list(a = 1, b = 1:3, c = 10:100)
lapply(x, FUN = length)
$a
[1] 1
$b
[1] 3
$c
[1] 91

lapply(x, FUN = sum)
$a
[1] 1
$b
[1] 6
$c
[1] 5005
```

- **sapply** - *When you want to apply a function to each element of a list in turn, but you want a **vector** back, rather than a list.*

If you find yourself typing `unlist(lapply(...))`, stop and consider `sapply`.

```
x <- list(a = 1, b = 1:3, c = 10:100)
#Compare with above; a named vector, not a list
sapply(x, FUN = length)
a b c
1 3 91

sapply(x, FUN = sum)
a b c
1 6 5005
```

In more advanced uses of `sapply` it will attempt to coerce the result to a multi-dimensional array, if appropriate. For example, if our function returns vectors of the same length, `sapply` will use them as columns of a matrix:

```
sapply(1:5,function(x) rnorm(3,x))
```

If our function returns a 2 dimensional matrix, `sapply` will do essentially the same thing, treating each returned matrix as a single long vector:

```
sapply(1:5,function(x) matrix(x,2,2))
```

Unless we specify `simplify = "array"`, in which case it will use the individual matrices to build a multi-dimensional array:

```
sapply(1:5,function(x) matrix(x,2,2), simplify = "array")
```

Each of these behaviors is of course contingent on our function returning vectors or matrices of the same length or dimension.

- **vapply** - *When you want to use `sapply` but perhaps need to squeeze some more speed out of your code.*

For `vapply`, you basically give R an example of what sort of thing your function will return, which can save some time coercing returned values to fit in a single atomic vector.

```
x <- list(a = 1, b = 1:3, c = 10:100)
#Note that since the advantage here is mainly speed, this
# example is only for illustration. We're telling R that
# everything returned by length() should be an integer of
# length 1.
vapply(x, FUN = length, FUN.VALUE = 0L)
a b c
1 3 91
```

- **mapply** - *For when you have several data structures (e.g. vectors, lists) and you want to apply a function to the 1st elements of each, and then the 2nd elements of each, etc., coercing the result to a vector/array as in `sapply`.*

This is multivariate in the sense that your function must accept multiple arguments.

```
#Sums the 1st elements, the 2nd elements, etc.
mapply(sum, 1:5, 1:5, 1:5)
[1] 3 6 9 12 15
#To do rep(1,4), rep(2,3), etc.
mapply(rep, 1:4, 4:1)
[[1]]
[1] 1 1 1 1

[[2]]
[1] 2 2 2

[[3]]
[1] 3 3

[[4]]
[1] 4
```

- **Map** - *A wrapper to `mapply` with `SIMPLIFY = FALSE`, so it is guaranteed to return a list.*

- `Map(sum, 1:5, 1:5, 1:5)`
- `[[1]]`
- `[1] 3`
- 
- `[[2]]`
- `[1] 6`
- 
- `[[3]]`
- `[1] 9`
- 
- `[[4]]`
- `[1] 12`
- 
- `[[5]]`
- `[1] 15`

- **rapply** - *For when you want to apply a function to each element of a **nested list** structure, recursively.*

To give you some idea of how uncommon `rapply` is, I forgot about it when first posting this answer! Obviously, I'm sure many people use it, but YMMV. `rapply` is best illustrated with a user-defined function to apply:

```
#Append ! to string, otherwise increment
myFun <- function(x){
  if (is.character(x)){
    return(paste(x,"!",sep=""))
  }
  else{
    return(x + 1)
  }
}

#A nested list structure
l <- list(a = list(a1 = "Boo", b1 = 2, c1 = "Eeek"),
         b = 3, c = "Yikes",
         d = list(a2 = 1, b2 = list(a3 = "Hey", b3 = 5)))

#Result is named vector, coerced to character
rapply(l,myFun)

#Result is a nested list like l, with values altered
rapply(l, myFun, how = "replace")
```

- **tapply** - *For when you want to apply a function to **subsets** of a vector and the subsets are defined by some other vector, usually a factor.*

The black sheep of the `*apply` family, of sorts. The help file's use of the phrase "ragged array" can be a bit [confusing](#), but it is actually quite simple.

A vector:

```
x <- 1:20
```

A factor (of the same length!) defining groups:

```
y <- factor(rep(letters[1:5], each = 4))
```

Add up the values in `x` within each subgroup defined by `y`:

```
tapply(x, y, sum)
  a  b  c  d  e
10 26 42 58 74
```

More complex examples can be handled where the subgroups are defined by the unique combinations of a list of several factors. `tapply` is similar in spirit to the split-apply-combine functions that are common in R (aggregate, by, ave, ddply, etc.) Hence its black sheep status.